

```
In [87]: import pandas as pd
import numpy as np

from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.metrics import precision_score, recall_score
```

Reading excel file

```
In [88]: data_file = pd.read_csv("Data_for_UCI_named.csv") # read csv file

In [89]: data_file

Out[89]:
```

	tau1	tau2	tau3	tau4	p1	p2	p3	p4	g1	g2	g3	g4	stabf
0	2.959060	3.079885	8.381025	9.780754	3.763085	-0.728204	-1.257395	-1.723086	0.850456	0.859578	0.887445	0.958034	0.05534
1	9.304097	4.902524	3.047541	1.369357	5.067812	-1.940058	-1.872742	-1.255012	0.413441	0.862414	0.562139	0.781760	-0.00594
2	8.971707	8.848428	3.046479	1.214518	3.405158	-1.207456	-1.277210	-0.920492	0.163041	0.766889	0.839444	0.109853	0.03047
3	0.716415	7.669600	4.486641	2.340563	3.963791	-1.027473	-1.938944	-0.997374	0.446209	0.976744	0.923931	0.362718	0.02867
4	3.134112	7.608772	4.943759	9.857573	3.525811	-1.125531	-1.845975	-0.554305	0.797110	0.455450	0.856947	0.820293	0.04986
...
9995	2.934006	9.487627	2.376524	6.187797	3.343416	-0.658054	-1.449106	-1.236256	0.601709	0.779642	0.813512	0.686385	0.02386
9996	3.392299	1.274827	2.954947	6.894759	4.349512	-1.663661	-0.952437	-1.733414	0.502079	0.567242	0.285880	0.366120	-0.02586
9997	2.364034	2.842030	8.776391	1.008906	4.299976	-1.380719	-0.943884	-1.975373	0.487838	0.986505	0.848918	0.145984	-0.03187
9998	9.631511	3.994393	2.750771	7.821347	2.514755	-0.966330	-0.649915	-0.898510	0.365246	0.587558	0.889118	0.818391	0.03776
9999	6.530527	6.781790	4.349695	8.673138	3.492807	-1.390285	-1.532193	-0.570329	0.073056	0.505441	0.378761	0.942631	0.04526

10000 rows × 14 columns

```
In [90]: data_file.columns

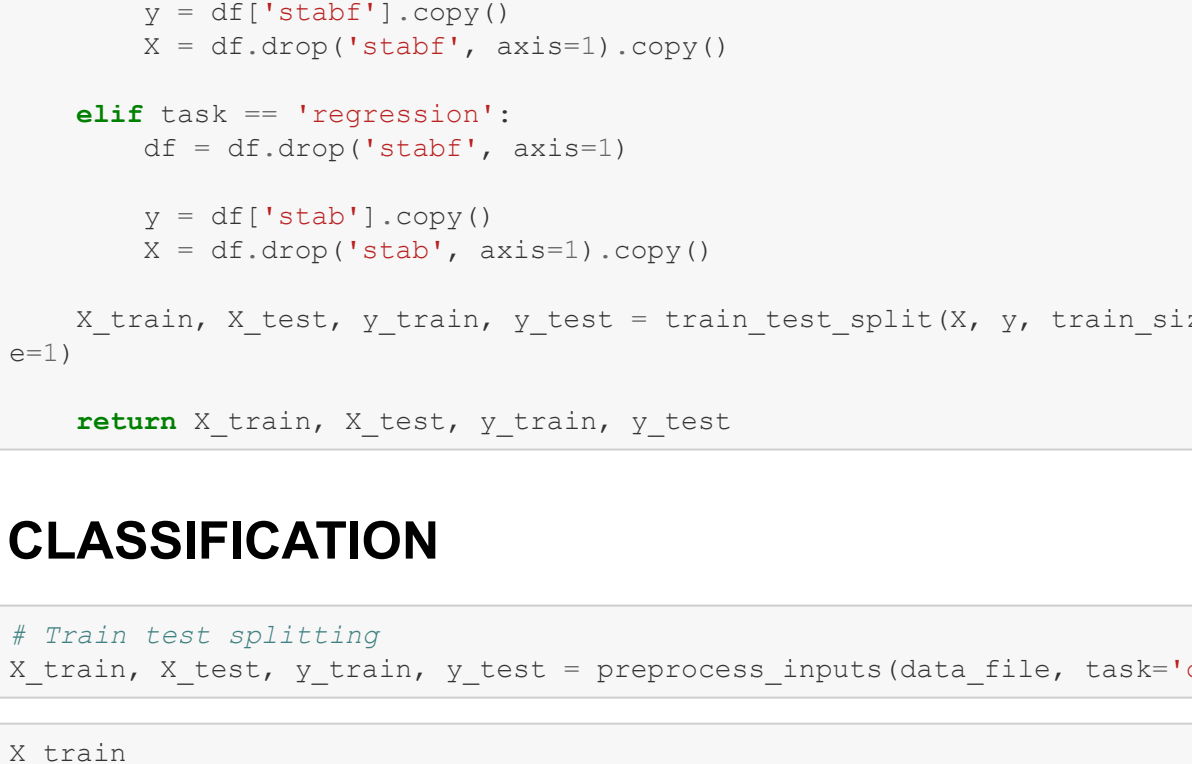
Out[90]: Index(['tau1', 'tau2', 'tau3', 'tau4', 'p1', 'p2', 'p3', 'p4', 'g1', 'g2', 'g3', 'g4', 'stabf', 'dtype='object'])

In [91]: data_file.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
tau1      10000 non-null float64
tau2      10000 non-null float64
tau3      10000 non-null float64
tau4      10000 non-null float64
p1         10000 non-null float64
p2         10000 non-null float64
p3         10000 non-null float64
p4         10000 non-null float64
g1         10000 non-null float64
g2         10000 non-null float64
g3         10000 non-null float64
g4         10000 non-null float64
stabf      10000 non-null object
dtypes: float64(13), object(1)
memory usage: 1.1+ MB
```

```
In [92]: import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline

In [93]: figure = plt.figure(figsize = (10,10))
column_list =list(data_file.columns[1:-2])
for i, v in enumerate(column_list):
    figure.add_subplot(3, 4, i+1)
    sns.distplot(data_file[v], hist = False) #plotting all the values
```



```
In [95]: # Separated the target and data from each other
# target = stabf and data is all the data except stabf and stab
def preprocess_inputs(df, task='classification'):
    df = df.copy()

    if task == 'classification':
        df = df.drop('stabf', axis=1)

        y = df['stabf'].copy()
        X = df.drop('stabf', axis=1).copy()

    elif task == 'regression':
        df = df.drop('stabf', axis=1)

        y = df['stabf'].copy()
        X = df.drop('stabf', axis=1).copy()

    X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=0.8, shuffle=True, random_state=1)

    return X_train, X_test, y_train, y_test

In [96]: # Train test splitting
X_train, X_test, y_train, y_test = preprocess_inputs(data_file, task='classification')

In [97]: X_train

Out[97]:
```

	tau1	tau2	tau3	tau4	p1	p2	p3	p4	g1	g2	g3	g4
2694	6.255995	2.542401	7.024714	9.476518	3.529888	-1.224881	-0.688228	-1.616780	0.568221	0.618403	0.685739	0.660088
5140	5.870581	5.490253	8.076688	0.761075	4.220888	-1.280596	-1.902185	-1.038107	0.483515	0.097244	0.916955	1.129254
2568	1.220702	8.804028	3.874283	8.433949	3.614027	-1.039236	-0.953566	-1.621224	0.908353	0.923594	0.238881	0.660156
3671	7.498402	6.697603	8.798626	2.126236	3.134585	-1.581906	-0.589386	-0.963293	0.260626	0.899003	0.964752	0.600598
7427	7.074006	1.337511	6.100756	7.591956	2.526922	-0.922540	-0.632600	-0.971782	0.984580	0.716082	0.836928	1.165162
...
2895	9.504984	5.265415	2.029617	8.026334	3.449756	-0.922754	-1.526305	-1.000696	0.191864	0.611297	0.080008	0.350312
7813	8.035818	4.632548	1.177810	1.204950	4.093208	-1.568218	-0.819966	-1.705024	0.223384	0.317770	0.726623	0.441620
905	7.052560	3.264057	1.405949	4.497549	3.741358	-0.629026	-1.407945	-1.704387	0.868306	0.605566	0.100591	0.747320
5192	5.085242	1.790031	2.486392	2.848192	3.159670	-0.892898	-1.025461	-1.241311	0.473917	0.418512	0.326253	0.962488
235	1.205354	6.995045	5.930860	2.523856	5.230127	-1.755047	-1.567927	-1.907153	0.861891	0.202768	0.848373	0.739993

8000 rows × 12 columns

```
In [98]: y_train

Out[98]:
```

2694	unstable
5140	unstable
2568	unstable
3671	unstable
7427	unstable
...	...
2895	stable
7813	stable
905	unstable
5192	unstable
235	stable

Name: stabf, Length: 8000, dtype: object

Stochastic Gradient Descent

```
In [99]: from sklearn.linear_model import SGDClassifier
sgd_clf = SGDClassifier()
sgd_clf.fit(X_train, y_train) # fitting data

Out[99]: SGDClassifier(alpha=0.0001, average=False, class_weight=None,
early_stopping=False, epsilon=0.1, eta0=0.0, fit_intercept=True,
l1_ratio=0.15, learning_rate='optimal', loss='hinge',
max_iter=1000, n_iter_no_change=5, n_jobs=None, penalty='l2',
power_t=0.5, random_state=None, shuffle=True, tol=0.001,
validation_fraction=0.1, verbose=0, warm_start=False)

In [100]: X_test[:5] # Checking for first 5 test data

Out[100]:
```

	tau1	tau2	tau3	tau4	p1	p2	p3	p4	g1	g2	g3	g4
9953	5.877876	4.113820	9.356768	8.299753	4.056779	-1.897470	-1.590581	-0.568728	0.276567	0.845536	0.112440	0.822562
3850	5.802841	6.271371	4.731540	3.819867	3.579569	-1.709480	-1.067511	-0.802579	0.077527	0.416478	0.912846	0.861306
4962	2.286998	4.385142	2.830232	5.293880	3.035814	-1.202764	-0.902011	-0.931039	0.924216	0.130186	0.703887	0.063811
3886	5.019920	2.209962	6.266080	0.578901	4.322584	-1.960207	-1.074561	-1.287815	0.546910	0.065992	0.427349	0.814648
5437	7.646145	9.187896	5.484219	9.934313	3.634226	-1.254541	-1.335366	-1.044319	0.561528	0.121611	0.787318	0.300314

```
In [101]: y_test # Checking for first 5 test data labels

Out[101]:
```

9953	unstable
3850	unstable
4962	stable
3886	stable
5437	unstable
...	...
3919	stable
162	stable
7903	stable
2242	unstable
2745	unstable

Name: stabf, Length: 2000, dtype: object

```
In [102]: sgd_clf.predict(X_test) # Predicting first 5 test data labels

Out[102]: array(['unstable', 'unstable', 'stable', ..., 'unstable', 'stable',
'unstable'], dtype='<U8')

In [103]: from sklearn.model_selection import cross_val_score
score = cross_val_score(sgd_clf, X_train, y_train, scoring='roc_auc', cv=4)
score.mean()

Out[103]: 0.881347713758724
```

```
In [104]: y_pred = sgd_clf.predict(X_test)

In [105]: confusion_matrix(y_test, y_pred) # confusion matrix values

Out[105]: array([[ 539, 1036],
[ 252, 1761]], dtype=int64)
```

```
In [106]: # precision score and recall score
precision_score(y_test, y_pred, pos_label="stable"), recall_score(y_test, y_pred, pos_label="stable")

Out[106]: (0.6814159292035398, 0.7570224719101124)
```

```
In [107]: print(classification_report(y_test, y_pred))

              precision    recall  f1-score   support

   stable         0.68         0.76         0.72         712
  unstable         0.86         0.80         0.83        1288

 accuracy          0.77         0.79         0.79        2000
 macro avg         0.77         0.78         0.77        2000
 weighted avg         0.79         0.79         0.79        2000
```

Support Vector Machines

```
In [108]: from sklearn import svm
svm_clf = svm.SVC()
svm_clf.fit(X_train, y_train)

C:\Users\17328\Anaconda3\lib\site-packages\sklearn\svm\base.py:193: FutureWarning: The default value
of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features.
Set gamma explicitly to 'auto' or 'scale' to avoid this warning.
"avoid this warning.", FutureWarning)

Out[108]: SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
decision_function_shape='ovr', degree=3, gamma='auto_deprecated',
kernel='rbf', max_iter=-1, probability=False, random_state=None,
shrinking=True, tol=0.001, verbose=False)

In [109]: X_test[:5] # Checking for first 5 test data

Out[109]:
```

	tau1	tau2	tau3	tau4	p1	p2	p3	p4	g1	g2	g3	g4
9953	6.877876	4.113820	9.356768	8.299753	4.056779	-1.897470	-1.590581	-0.568728	0.276567	0.845536	0.112440	0.822562
3850	5.802841	6.271371	4.731540	3.819867	3.579569	-1.709480	-1.067511	-0.802579	0.077527	0.416478	0.912846	0.861306
4962	2.286998	4.385142	2.830232	5.293880	3.035814	-1.202764	-0.902011	-0.931039	0.924216	0.130186	0.703887	0.063811
3886	5.019920	2.209962	6.266080	0.578901	4.322584	-1.960207	-1.074561	-1.287815	0.546910	0.065992	0.427349	0.814648
5437	7.646145	9.187896	5.484219	9.934313	3.634226	-1.254541	-1.335366	-1.044319	0.561528	0.121611	0.787318	0.300314

```
In [110]: y_test[:5] # Checking for first 5 test data

Out[110]:
```

9953	unstable
3850	unstable
4962	stable
3886	stable
5437	unstable

Name: stabf, dtype: object

```
In [111]: svm_clf.predict(X_test[:5]) # Predicting first 5 test data

Out[111]: array(['unstable', 'unstable', 'stable', 'stable', 'unstable'],
dtype=object)
```

```
In [112]: y_pred_svm = svm_clf.predict(X_test) # predicting all test data

In [113]: confusion_matrix(y_test, y_pred_svm)

Out[113]: array([[ 602, 110],
[ 61, 1227]], dtype=int64)
```

```
In [114]: precision_score(y_test, y_pred_svm, pos_label = "stable"), recall_score(y_test, y_pred_svm, pos_label = "stable")

Out[114]: (0.9079939668174962, 0.8455056179775281)
```

```
In [115]: print(classification_report(y_test, y_pred_svm))

              precision    recall  f1-score   support

   stable         0.91         0.85         0.88         712
  unstable         0.92         0.95         0.93        1288

 accuracy          0.91         0.91         0.91        2000
 macro avg         0.91         0.90         0.91        2000
 weighted avg         0.91         0.91         0.91        2000
```

Naive Bayes

```
In [116]: from sklearn.naive_bayes import GaussianNB

In [117]: gnb = GaussianNB()

In [118]: gnb.fit(X_train, y_train)

Out[118]: GaussianNB(priors=None, var_smoothing=1e-09)

In [119]: X_test[:5] # Checking for first 5 test data

Out[119]:
```

	tau1	tau2	tau3	tau4	p1	p2	p3	p4	g1	g2	g3	g4
9953	6.877876	4.113820	9.356768	8.299753	4.056779	-1.897470	-1.590581	-0.568728	0.276567	0.845536	0.112440	0.822562
3850	5.802841	6.271371	4.731540	3.819867	3.579569	-1.709480	-1.067511	-0.802579	0.077527	0.416478	0.912846	0.861306
4962	2.286998	4.385142	2.830232	5.293880	3.035814	-1.202764	-0.902011	-0.931039	0.924216	0.130186	0.703887	0.063811
3886	5.019920	2.209962	6.266080	0.578901	4.322584	-1.960207	-1.074561	-1.287815	0.546910	0.065992	0.427349	0.814648
5437	7.646145	9.187896	5.484219	9.934313	3.634226	-1.254541	-1.335366	-1.044319	0.561528	0.121611	0.787318	0.300314

```
In [120]: y_test[:5] # Checking for first 5 test data

Out[120]:
```

9953	unstable
3850	unstable
4962	stable
3886	stable
5437	unstable

Name: stabf, dtype: object

```
In [121]: gnb.predict(X_test[:5]) # Predicting first 5 test data

Out[121]: array(['unstable', 'unstable', 'stable', 'stable', 'unstable'],
dtype='<U8')

In [122]: y_pred_gnb = gnb.predict(X_test) # Predicting all test data

In [123]: confusion_matrix(y_test, y_pred_gnb)

Out[123]: array([[ 476, 236],
[ 106, 1182]], dtype=int64)
```

```
In [124]: precision_score(y_test, y_pred_gnb, pos_label = "stable"), recall_score(y_test, y_pred_gnb, pos_label = "stable")

Out[124]: (0.8178694158075601, 0.6685393258426966)
```

```
In [125]: print(classification_report(y_test, y_pred_gnb))

              precision    recall  f1-score   support

   stable         0.82         0.67         0.74         712
  unstable         0.83         0.92         0.87        1288

 accuracy          0.83         0.79         0.83        2000
 macro avg         0.83         0.79         0.80        2000
 weighted avg         0.83         0.83         0.82        2000
```

Decision Trees

```
In [126]: from sklearn import tree
tree_clf = tree.DecisionTreeClassifier()
tree_clf.fit(X_train, y_train)

Out[126]: DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=None,
max_features=None, max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, presort=False,
random_state=None, splitter='best')

In [127]: X_test[:5] # Checking for first 5 test data

Out[127]:
```

	tau1	tau2	tau3	tau4	p1	p2	p3	p4	g1	g2	g3	g4
9953	6.877876	4.113820	9.356768	8.299753	4.056779	-1.897470	-1.590581	-0.568728	0.276567	0.845536	0.112440	0.822562
3850	5.802841	6.271371	4.731540	3.819867	3.579569	-1.709480	-1.067511	-0.802579	0.077527	0.416478	0.912846	0.861306
4962	2.286998	4.385142	2.830232	5.293880	3.035814	-1.202764	-0.902011	-0.931039	0.924216	0.130186	0.703887	0.063811
3886	5.019920	2.209962	6.266080	0.578901	4.322584	-1.960207	-1.074561	-1.287815	0.546910	0.065992	0.427349	0.814648
5437	7.646145	9.187896	5.484219	9.934313	3.634226	-1.254541	-1.335366	-1.044319	0.561528	0.121611	0.787318	0.300314

```
In [128]: y_test[:5] # Checking for first 5 test data

Out[128]:
```

9953	unstable
3850	unstable
4962	stable
3886	stable
5437	unstable

Name: stabf, dtype: object

```
In [129]: tree_clf.predict(X_test[:5]) # Predicting first 5 test data

Out[129]: array(['unstable', 'unstable', 'stable', 'stable', 'unstable'],
dtype=object)
```

```
In [130]: y
```